

# A Tutorial on Inference and Learning in Bayesian Networks

**Irina Rish**

IBM T.J.Watson Research Center

[rish@us.ibm.com](mailto:rish@us.ibm.com)

<http://www.research.ibm.com/people/r/rish/>

# Outline

- Motivation: learning probabilistic models from data
- Representation: Bayesian network models
- Probabilistic inference in Bayesian Networks
  - Exact inference
  - Approximate inference
- Learning Bayesian Networks
  - Learning parameters
  - Learning graph structure (model selection)
- Summary

# Bayesian Networks

Structured, graphical representation of probabilistic relationships between several random variables

Explicit representation of conditional independencies

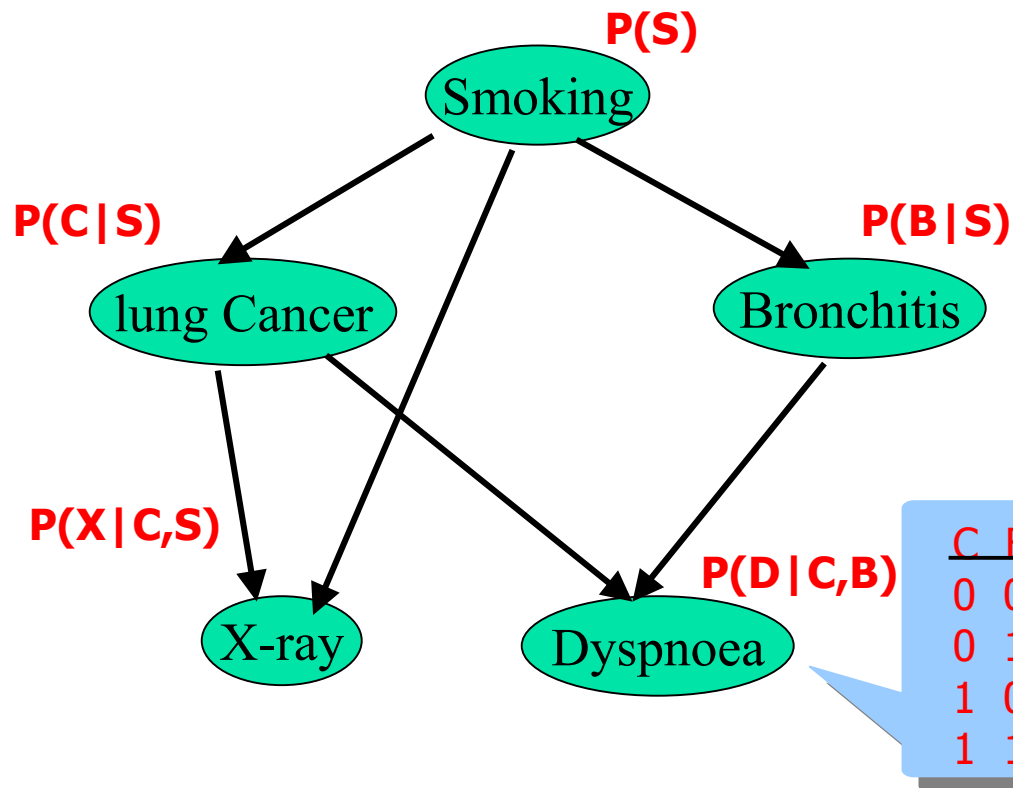
Missing arcs encode conditional independence

Efficient representation of joint PDF  $P(X)$

Generative model (not just discriminative): allows arbitrary queries to be answered, e.g.

$P(\text{lung cancer}=\text{yes} \mid \text{smoking}=\text{no}, \text{positive X-ray}=\text{yes}) = ?$

# Bayesian Network: $\mathbf{BN} = (\mathbf{G}, \Theta)$



$\mathbf{G}$  - directed acyclic graph (DAG)  
nodes – random variables  
edges – direct dependencies

$\Theta$  - set of parameters in all conditional probability distributions (CPDs)

CPD:

| C | B | D=0 | D=1 |
|---|---|-----|-----|
| 0 | 0 | 0.1 | 0.9 |
| 0 | 1 | 0.7 | 0.3 |
| 1 | 0 | 0.8 | 0.2 |
| 1 | 1 | 0.9 | 0.1 |

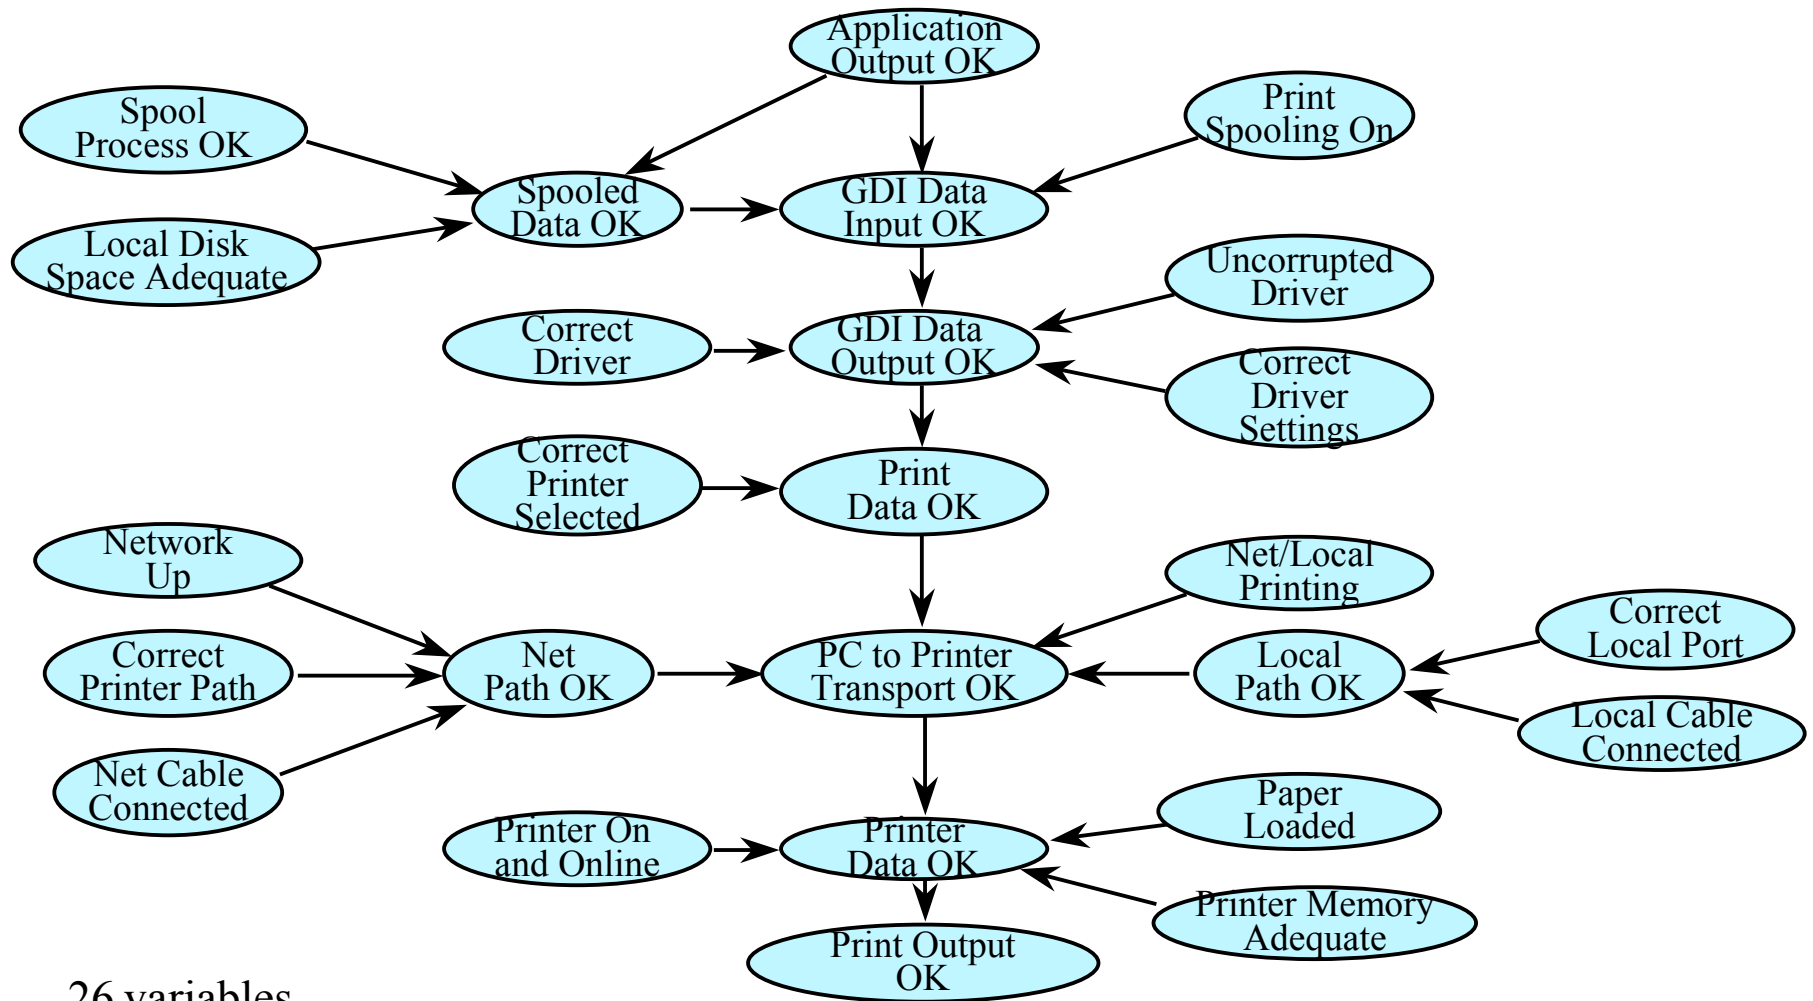
**CPD of  
node X:  
 $P(X | \text{parents}(X))$**

**Compact representation** of joint distribution in a **product form** (chain rule):

$$P(S, C, B, X, D) = P(S) P(C|S) P(B|S) P(X|C,S) P(D|C,B)$$

$$1 + 2 + 2 + 4 + 4 = 13 \text{ parameters instead of } 2^5 = 32$$

# Example: Printer Troubleshooting



26 variables

Instead of  $2^{26}$  parameters we get

$$99 = 17x1 + 1x2^1 + 2x2^2 + 3x2^3 + 3x2^4$$

# "Moral" graph of a BN

Moralization algorithm:


1. Connect ("marry") parents of each node.
2. Drop the directionality of the edges.

Resulting undirected graph is called the "moral" graph of BN

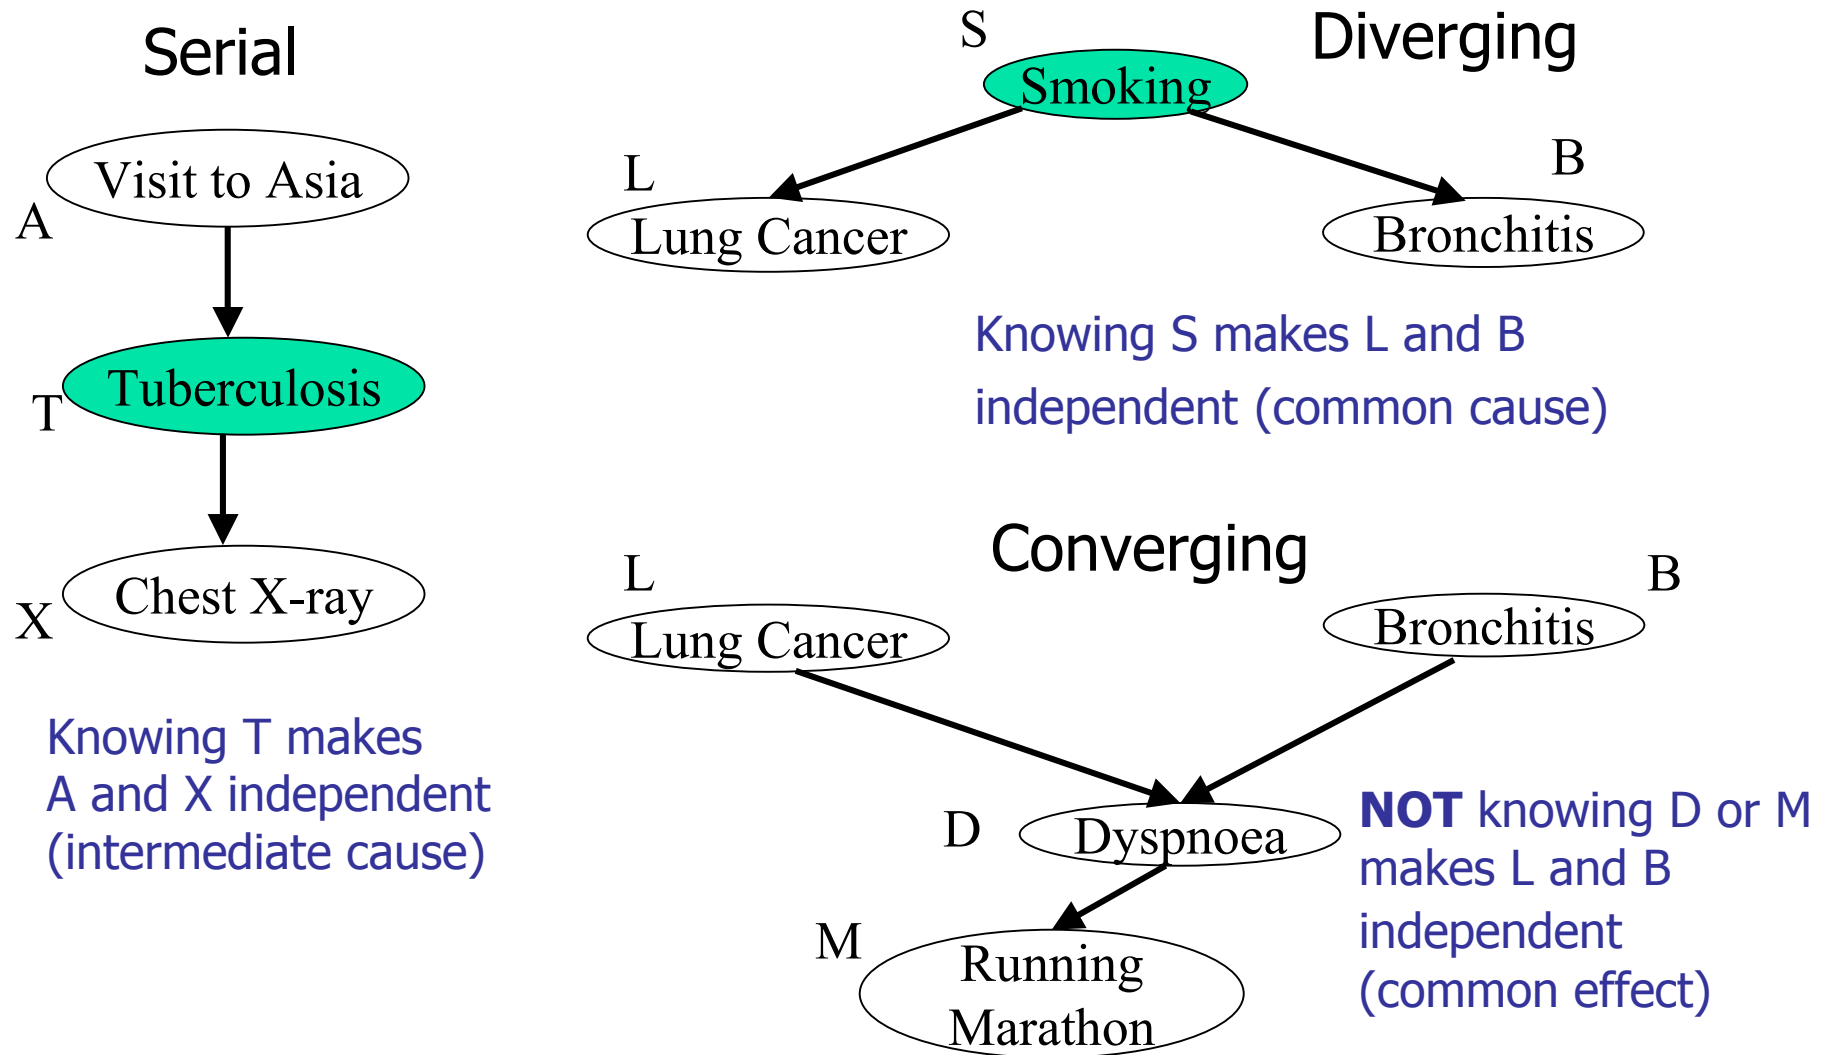
Interpretation:

every pair of nodes that occur together in a CPD is connected by an edge in the moral graph.

CPD for  $X$  and its  $k$  parents (called "**family**") is represented by a clique of size **( $k+1$ )** in the moral graph, and contains  $d^k (d-1)$  probability parameters where  $d$  is the number of values each variable can have (domain size).



# Conditional Independence in BNs: Three types of connections




# d-separation

Nodes X and Y are *d-separated* if on *any (undirected) path* between X and Y there is some variable Z such that is either

Z is in a *serial* or *diverging* connection and Z is *known*, or

Z is in a *converging* connection and neither Z nor any of Z's descendants are known



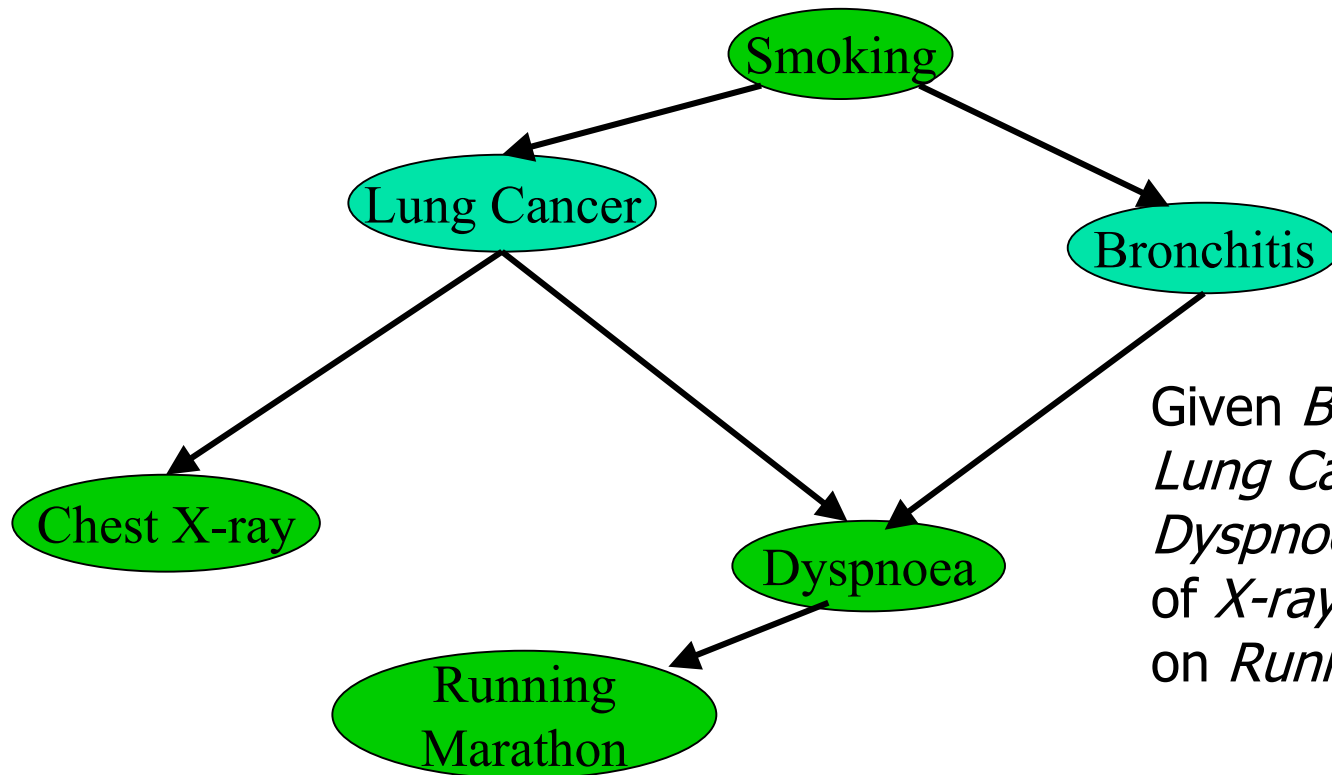
Nodes X and Y are called *d-connected* if they are not d-separated (there exists an undirected path between X and Y not d-separated by any node or a set of nodes)

If nodes X and Y are *d-separated* by Z, then X and Y are *conditionally independent* given Z (see Pearl, 1988)



# Independence Relations in BN

A variable (node) is conditionally independent of its non-descendants given its parents

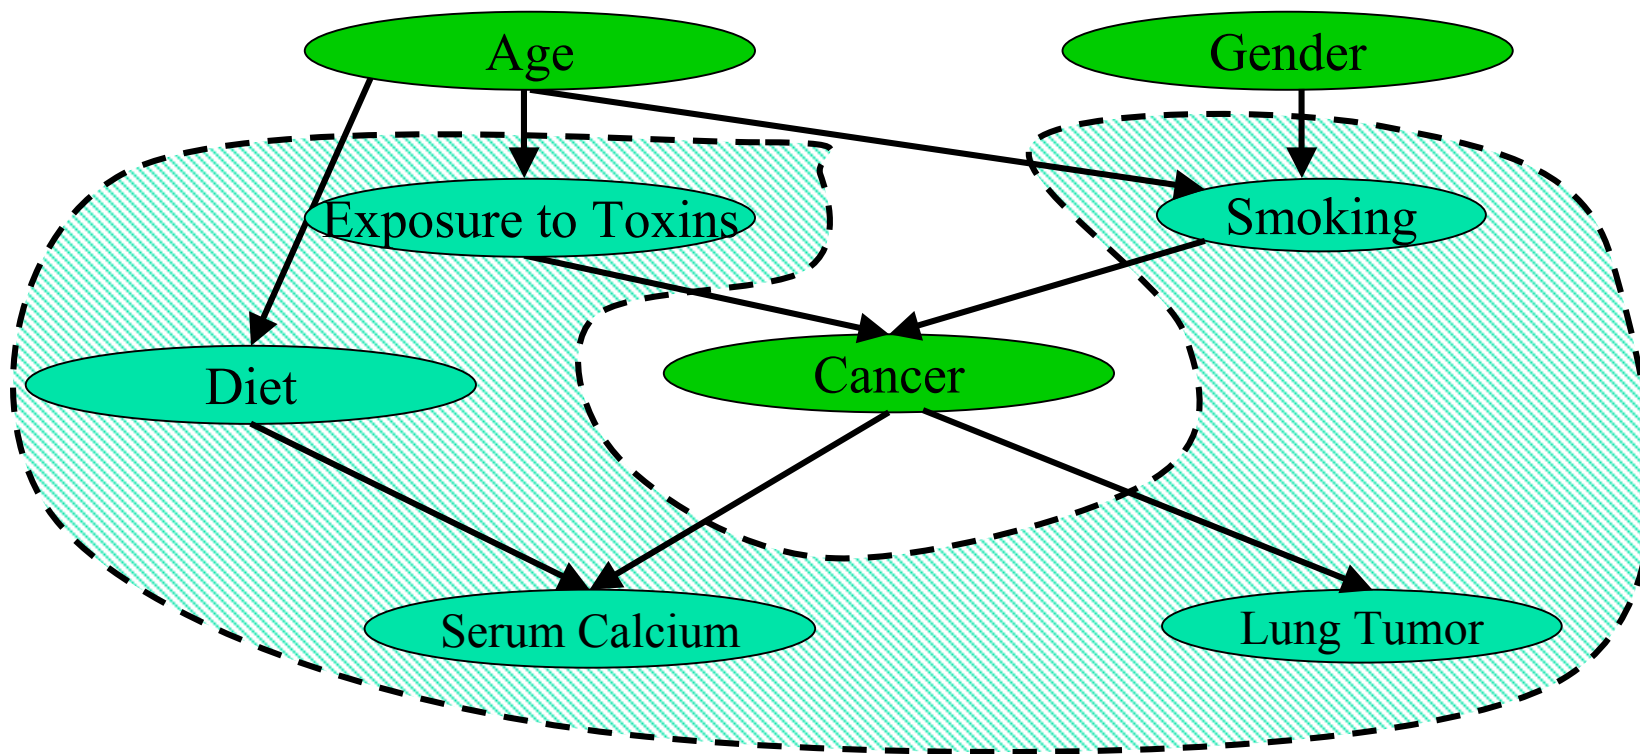


Given *Bronchitis* and *Lung Cancer*, *Dyspnoea* is independent of *X-ray* (but may depend on *Running Marathon*)

## Markov Blanket



A node is conditionally independent of ALL other nodes given its *Markov blanket*, i.e. its *parents*, *children*, and “*spouses*” (parents of common children)

(Proof left as a homework problem ☺)




# What are BNs useful for?

- Diagnosis:  $P(\text{cause}|\text{symptom})=?$
- Prediction:  $P(\text{symptom}|\text{cause})=?$
- Classification:  $\max_{\text{class}} P(\text{class}|\text{data})$
- Decision-making (given a cost function)



Speech  
recognition



# Application Examples

APRI system developed at AT&T Bell Labs

- learns & uses Bayesian networks from data to identify customers liable to default on bill payments

NASA Vista system

- predict failures in propulsion systems

- considers time criticality & suggests highest utility action

- dynamically decide what information to show

# Application Examples

## Office Assistant in MS Office 97/ MS Office 95

- Extension of Answer wizard

- uses naïve Bayesian networks

- help based on past experience (keyboard/mouse use) and task user is doing currently

- This is the “smiley face” you get in your MS Office applications

## Microsoft Pregnancy and Child-Care

- Available on MSN in Health section

- Frequently occurring children’s symptoms are linked to expert modules that repeatedly ask parents relevant questions

- Asks next best question based on provided information

- Presents articles that are deemed relevant based on information provided


# IBM's systems management applications

Machine Learning for Systems @ Watson

(Hellerstein, Jayram, Rish (2000))

## End-user transaction recognition


Remote Procedure Calls (RPCs)



(Rish, Brodie, Ma (2001))

## Fault diagnosis using probes

Software or hardware components



### Issues:

- Efficiency (scalability)
- Missing data/noise: sensitivity analysis
- "Adaptive" probing:
  - selecting "most-informative" probes
- on-line learning/model updates
- on-line diagnosis

**Goal: finding most-likely diagnosis**

$$(x_1^*, \dots, x_n^*) = \arg \max_{x_1, \dots, x_n} P(x_1, \dots, x_n | t_1, \dots, t_n)$$

**Pattern discovery, classification, diagnosis and prediction**

# Probabilistic Inference Tasks

- Belief updating:

$$\text{BEL}(X_i) = P(X_i = x_i \mid \text{evidence})$$

- Finding most probable explanation (MPE)

$$\bar{x}^* = \operatorname{argmax}_{\bar{x}} P(\bar{x}, e)$$


- Finding maximum a-posteriori hypothesis

$$(a_1^*, \dots, a_k^*) = \operatorname{argmax}_{\bar{a}} \sum_{\bar{x}/\bar{a}} P(\bar{x}, e) \quad \begin{array}{l} A \subseteq X: \\ \text{hypothesis variables} \end{array}$$

- Finding maximum-expected-utility (MEU) decision

$$(d_1^*, \dots, d_k^*) = \operatorname{argmax}_{\bar{d}} \sum_{\bar{x}/\bar{d}} P(\bar{x}, e) U(\bar{x}) \quad \begin{array}{l} D \subseteq X: \text{decision variables} \\ U(\bar{x}): \text{utility function} \end{array}$$


# Belief Updating Task: Example



$P(\text{smoking} \mid \text{dyspnoea}=\text{yes}) = ?$



# Belief updating: find $P(X|\text{evidence})$



$$P(s|d=1) = \frac{P(s, d=1)}{P(d=1)} \propto P(s, d=1) =$$

$$\sum_{d=1, b, x, c} P(s) \underbrace{P(c|s)} P(b|s) \underbrace{P(x|c, s) P(d|c, b)} =$$

$$P(s) \sum_{d=1} \sum_b P(b|s) \sum_x \underbrace{\sum_c P(c|s) P(x|c, s) P(d|c, b)}_{f(s, d, b, x)}$$

$W^*=4$

Variable Elimination

Complexity:  $O(n \exp(w^*))$

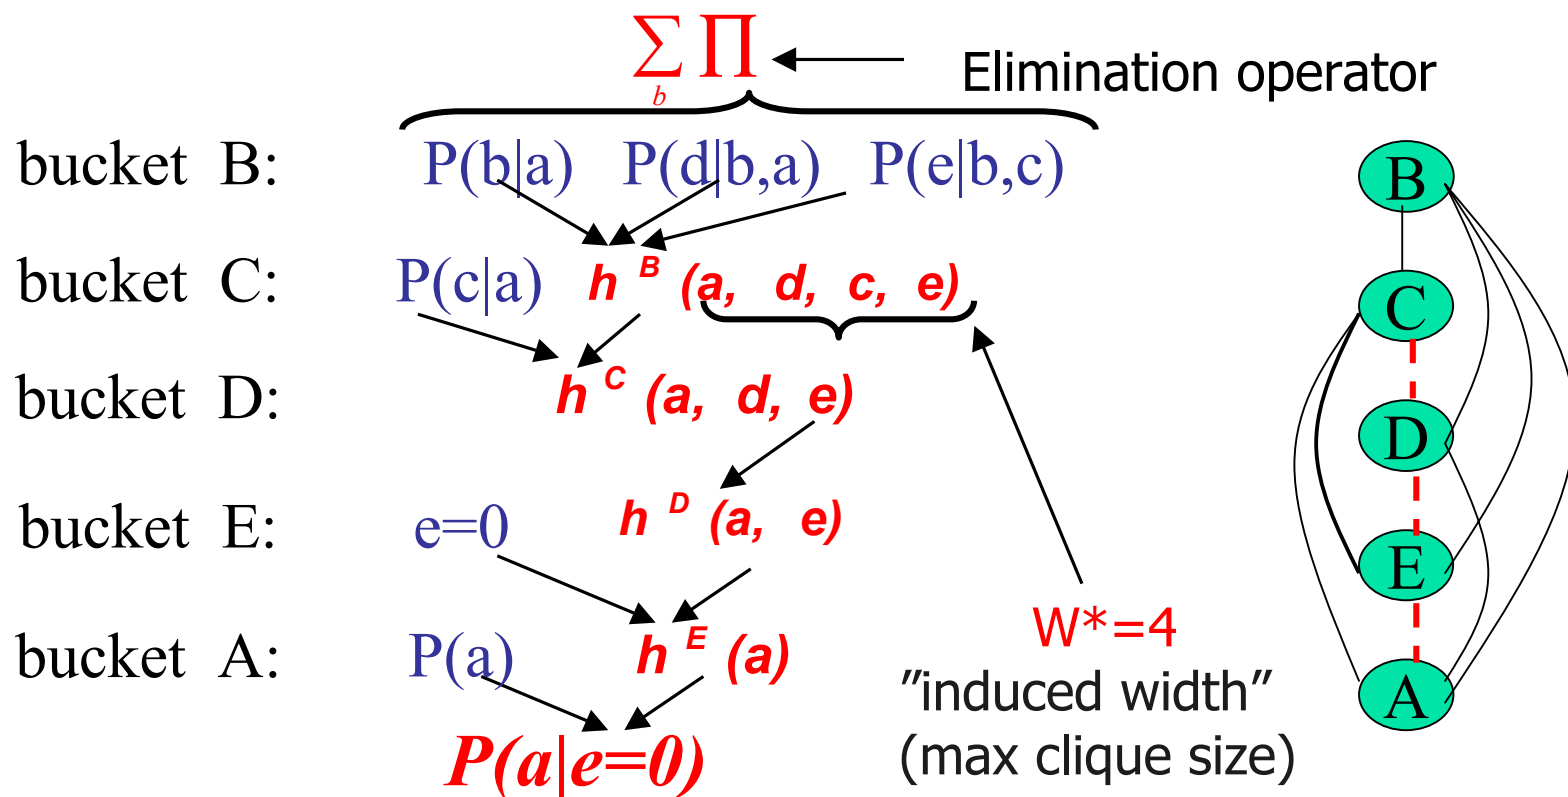
"induced width"  
(max induced clique size)

Efficient inference: variable orderings, conditioning, approximations

# Variable elimination algorithms

(also called "bucket-elimination")

Belief updating: **VE-algorithm** *elim-bel* (Dechter 1996)




# Finding $MPE = \max_{\bar{x}} P(\bar{x})$

VE-algorithm *elim-mpe* (Dechter 1996)

$$\sum \text{ is replaced by } \max :$$

$$MPE = \max_{a,e,d,c,b} P(a)P(c|a)P(b|a)P(d|a,b)P(e|b,c)$$



# Generating the MPE-solution

$$5. \quad b' = \arg \max_b P(b \mid a') \times \\ \times P(d' \mid b, a') \times P(e' \mid b, c')$$

$$4. \quad c' = \arg \max_c P(c \mid a') \times \\ \times h^B(a', d', c, e')$$

$$3. \quad d' = \arg \max_d h^C(a', d, e')$$

$$2. \quad e' = 0$$

$$1. \quad a' = \arg \max_a P(a) \cdot h^E(a)$$

$$B: \quad P(b|a) \quad P(d|b,a) \quad P(e|b,c)$$

$$C: \quad P(c|a) \quad h^B(a, d, c, e)$$

$$D: \quad h^C(a, d, e)$$

$$E: \quad e=0 \quad h^D(a, e)$$

$$A: \quad P(a) \quad h^E(a)$$

Return  $(a', b', c', d', e')$

# Complexity of VE-inference: $O(n \exp(w_o^*))$

$w_o^*$  – the induced width of moral graph along ordering  $o$


Meaning :  $w_o^* + 1 =$  size of largest clique created during inference

The width  $w_o(X)$  of a variable  $X$  in graph  $G$  along the ordering  $o$  is the number of nodes preceding  $X$  in the ordering and connected to  $X$  (*earlier neighbors*).


The width  $w_o$  of a graph is the maximum width  $w_o(X)$  among all nodes.

The *induced graph*  $G'$  along the ordering  $o$  is obtained by recursively connecting earlier neighbors of each node, from last to the first in the ordering.


The width of the induced graph  $G'$  is called the *induced width* of the graph  $G$  (denoted  $w_o^*$ ).



"Moral" graph



$$w_{o_1}^* = 4$$




$$w_{o_2}^* = 2$$


**Ordering is important! But finding min- $w^*$  ordering is NP- hard... Inference is also NP-hard in general case [Cooper].**

# Learning Bayesian Networks

- Combining domain expert knowledge with data



- Efficient representation and inference



- Incremental learning:  $P(H) \nearrow$  or  $\searrow$


- Handling missing data: **<1.3 2.8 ?? 0 1>**

- Learning causal relationships: 

# Learning tasks: four main cases


## ■ Known graph – learn parameters

- Complete data:  
parameter estimation (ML, MAP)
- Incomplete data:  
non-linear parametric  
optimization (gradient descent, EM)



## ■ Unknown graph – learn graph and parameters

- Complete data:  
optimization (search  
in space of graphs)
- Incomplete data:  
structural EM,  
mixture models




$$\hat{G} = \arg \max_G \text{Score}(G)$$

# Learning Parameters: complete data

## (overview)

- ML-estimate:  $\max_{\Theta} \underbrace{\log P(D | \Theta)}_{\text{decomposable!}}$



$$\text{ML}(\theta_{x, \text{pa}_X}) = \frac{N_{x, \text{pa}_X}}{\sum_x N_{x, \text{pa}_X}}$$

counts

- MAP-estimate  
(Bayesian statistics)

$$\max_{\Theta} \underbrace{\log P(D | \Theta) P(\Theta)}$$

**Conjugate** priors - **Dirichlet**  $\text{Dir}(\theta_{\text{pa}_X} | \alpha_{1, \text{pa}_X}, \dots, \alpha_{m, \text{pa}_X})$


$$\text{MAP}(\theta_{x, \text{pa}_X}) = \frac{N_{x, \text{pa}_X} + \alpha_{x, \text{pa}_X}}{\sum_x N_{x, \text{pa}_X} + \underbrace{\sum_x \alpha_{x, \text{pa}_X}}_{\text{Equivalent sample size (prior knowledge)}}}$$

Equivalent sample size  
(prior knowledge)



# Likelihood Function

◆ By definition of network, we get




$$L(\Theta : D) = \prod_m P(E[m], B[m], A[m], C[m] : \Theta)$$

$$= \prod_m \begin{pmatrix} P(E[m] : \Theta) \\ P(B[m] : \Theta) \\ P(A[m] | B[m], E[m] : \Theta) \\ P(C[m] | A[m] : \Theta) \end{pmatrix}$$



# Likelihood Function

◆ Rewriting terms, we get




$$L(\Theta : D) = \prod_m P(E[m], B[m], A[m], C[m] : \Theta)$$

$$\prod_m P(E[m] : \Theta)$$

$$\prod_m P(B[m] : \Theta)$$

$$= \prod_m P(A[m] | B[m], E[m] : \Theta)$$

$$\prod_m P(C[m] | A[m] : \Theta)$$



# General Bayesian Networks

Generalizing for any Bayesian network:

$$\begin{aligned} L(\Theta : D) &= \prod_m P(x_1[m], \dots, x_n[m] : \Theta) \\ &= \prod_i \prod_m P(x_i[m] \mid Pa_i[m] : \Theta_i) \\ &= \prod_i L_i(\Theta_i : D) \end{aligned}$$


**Decomposition**

**$\Rightarrow$  Independent estimation problems**

# Likelihood Function: Multinomials

$$L(\theta : D) = P(D | \theta) = \prod_m P(x[m] | \theta)$$

◆ The likelihood for the sequence H, T, T, H, H is



$$L(\theta : D) = \theta \cdot (1 - \theta) \cdot (1 - \theta) \cdot \theta \cdot \theta$$

General case:

$$L(\Theta : D) = \prod_{k=1}^K \theta_k^{N_k}$$

Count of  $k^{\text{th}}$   
outcome in  $D$

Probability of  
 $k^{\text{th}}$  outcome

# Dirichlet Priors

◆ Recall that the likelihood function is

$$L(\Theta : D) = \prod_{k=1}^K \theta_k^{N_k}$$

◆ **Dirichlet** prior with hyperparameters  $\alpha_1, \dots, \alpha_K$

$$P(\Theta) \propto \prod_{k=1}^K \theta_k^{\alpha_k - 1}$$

$\Rightarrow$  the posterior has the same form, with

hyperparameters  $\alpha_1 + N_1, \dots, \alpha_K + N_K$

$$P(\Theta | D) \propto P(\Theta)P(D | \Theta) \propto \prod_{k=1}^K \theta_k^{\alpha_k - 1} \prod_{k=1}^K \theta_k^{N_k} = \prod_{k=1}^K \theta_k^{\alpha_k + N_k - 1}$$

# Learning Parameters: Summary

◆ Estimation relies on **sufficient statistics**

- For multinomials: counts  $N(x_i, pa_i)$
- Parameter estimation

$$\hat{\theta}_{x_i|pa_i} = \frac{N(x_i, pa_i)}{N(pa_i)} \quad \tilde{\theta}_{x_i|pa_i} = \frac{\alpha(x_i, pa_i) + N(x_i, pa_i)}{\alpha(pa_i) + N(pa_i)}$$

MLE Bayesian (Dirichlet)

- ◆ Both are asymptotically equivalent and consistent
- ◆ Both can be implemented in an on-line manner by accumulating sufficient statistics

# Learning Parameters: incomplete data

Non-decomposable marginal likelihood (hidden nodes)

Initial parameters

Current model  
( $G, \Theta$ )

Expectation

Compute EXPECTED  
Counts via inference in BN

Data

| S           | X | D | C | B |
|-------------|---|---|---|---|
| <? 0 1 0 1> |   |   |   |   |
| <1 1 ? 0 1> |   |   |   |   |
| <0 0 0 ? ?> |   |   |   |   |
| <? ? 0 ? 1> |   |   |   |   |

Expected counts

$$E_{P(X)}[N_{x, \text{pa}_x}] = \sum_{k=1}^N p(x, \text{pa}_x | y^k, \Theta, G)$$

Maximization

Update parameters  
(ML, MAP)

EM-algorithm:  
iterate until convergence

# Learning graph structure

$$\text{Find } \hat{\mathbf{G}} = \arg \max_G \text{Score}(\mathbf{G})$$

NP-hard  
optimization

## ■ Heuristic search:


Complete data – local computations

Incomplete data (score non-decomposable): stochastic methods

Local greedy search; K2 algorithm

## ■ Constrained-based methods (PC/IC algorithms)

➤ Data impose independence relations (constraints) on graph structure





# Scoring function:

## Minimum Description Length (MDL)


- Learning  $\Leftrightarrow$  data compression

$$MDL(BN \mid D) = \underbrace{-\log P(D \mid \Theta, G)}_{DL(\text{Data}|\text{model})} + \underbrace{\frac{\log N}{2} |\Theta|}_{DL(\text{Model})}$$


- Other: MDL = -BIC (Bayesian Information Criterion)
- Bayesian score (BDe) - asymptotically equivalent to MDL

# Model selection trade-offs

**Naïve Bayes** – too simple  
(less parameters, but bad model)



**Unrestricted BN** – too complex  
(possible overfitting + complexity)




**Various approximations between the two extremes**

## TAN:

tree-augmented Naïve Bayes  
[Friedman et al. 1997]

Based on Chow-Liu Tree Method  
(CL) for learning trees  
[Chow-Liu, 1968]




# Tree-structured distributions


A joint probability distribution is tree-structured if it can be written as

$$P(\mathbf{X}) = \prod_{i=1}^n P(x_i \mid x_{j(i)})$$

where  $x_{j(i)}$  is the parent of  $x_i$  in Bayesian network for  $P(\mathbf{x})$  (a directed tree)



A tree (with root A)




Not a tree – has an (undirected) cycle


A tree requires only  **$[(d-1) + d(d-1)(n-1)]$**  parameters, where  $d$  is domain size  
Moreover, inference in trees is  $O(n)$  (linear) since their  $w^*=1$

# Approximations by trees

True distribution  $P(X)$



Tree-approximation  $P'(X)$



How good is approximation? Use cross-entropy (KL-divergence):

$$D(P, P') = P \sum_{\mathbf{x}} P(\mathbf{x}) \log \frac{P(\mathbf{x})}{P'(\mathbf{x})}$$

$D(P, P')$  is non-negative, and  $D(P, P')=0$  if and only if  $P$  coincides with  $P'$  (on a set of measure 1)

How to find the best tree-approximation?

# Optimal trees: Chow-Liu result

- **Lemma**

Given a joint PDF  $P(x)$  and a fixed tree structure  $T$ , the best approximation  $P'(x)$  (i.e.,  $P'(x)$  that minimizes  $D(P, P')$ ) satisfies

$$P'(x_i | x_{j(i)}) = P(x_i | x_{j(i)}) \quad \text{for all } i = 1, \dots, n$$

Such  $P'(x)$  is called the projection of  $P(x)$  on  $T$ .

- **Theorem** [Chow and Liu, 1968]

Given a joint PDF  $P(x)$ , the KL-divergence  $D(P, P')$  is minimized by projecting  $P(x)$  on a *maximum-weight spanning tree (MSWT)* over nodes in  $X$ , where the weight on the edge  $(X_i, X_j)$  is defined by the mutual information measure

$$I(X_i; X_j) = \sum_{x_i, x_j} P(x_i, x_j) \log \frac{P(x_i, x_j)}{P(x_i)P(x_j)}$$

Note, that  $I(X; Y) = 0$  when  $X$  and  $Y$  are independent and that  $I(X; Y) = D(P(x, y), P(x)P(y))$

# Proofs

## Proof of Lemma :

$$\begin{aligned}
 D(P, P') &= -\sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1}^n \log P'(x_i | x_{j(i)}) + \sum_{\mathbf{x}} P(\mathbf{x}) \log P(\mathbf{x}) = -\sum_{\mathbf{x}} P(\mathbf{x}) \sum_{i=1}^n \log P'(x_i | x_{j(i)}) - H(X) = \\
 &= -\sum_{i=1}^n \sum_{\mathbf{x}} P(\mathbf{x}) \log P'(x_i | x_{j(i)}) - H(X) = -\sum_{i=1}^n \sum_{x_i, x_{j(i)}} P(x_i, x_{j(i)}) \log P'(x_i | x_{j(i)}) - H(X) = \quad (1)
 \end{aligned}$$

$$= -\sum_{i=1}^n \sum_{x_{j(i)}} P(x_{j(i)}) \sum_{x_i} P(x_i | x_{j(i)}) \log P'(x_i | x_{j(i)}) - H(X) \quad (2)$$

A known fact : given  $P(\mathbf{x})$ , the maximum of  $\sum_{x_i} P(x_i) \log P'(x_i)$  is achieved by the choice  $P'(x_i) = P(x_i)$ .

Therefore, for any value of  $i$  and  $x_{j(i)}$ , the term  $\sum_{x_i} P(x_i | x_{j(i)}) \log P'(x_i | x_{j(i)})$  is maximized by

choosing  $P'(x_i | x_{j(i)}) = P(x_i | x_{j(i)})$  (and thus the total  $D(P, P')$  is minimized), which proves the Lemma.

## Proof of Theorem :

Replacing  $P'(x_i | x_{j(i)}) = P(x_i | x_{j(i)})$  in the expression (1) yields

$$\begin{aligned}
 D(P, P') &= -\sum_{i=1}^n \sum_{x_i, x_{j(i)}} P(x_i, x_{j(i)}) \log [P(x_i x_{j(i)}) / P(x_{j(i)})] - H(X) = \\
 &= -\sum_{i=1}^n \sum_{x_i, x_{j(i)}} P(x_i, x_{j(i)}) \left[ \log \frac{P(x_i x_{j(i)})}{P(x_{j(i)}) P(x_i)} + \log P(x_i) \right] - H(X) = \\
 &= -\sum_{i=1}^n I(X_i, X_{j(i)}) - \sum_{i=1}^n \sum_{x_i} P(x_i) \log P(x_i) - H(X).
 \end{aligned}$$

The last two terms are independent of the choice of the tree, and thus  $D(P, P')$  is minimized by maximizing the sum of edge weights  $I(X_i, X_{j(i)})$ .

# Chow-Liu algorithm


[As presented in Pearl, 1988]

1. From the given distribution  $P(x)$  (or from data generated by  $P(x)$ ), compute the joint distribution  $P(x_i | x_j)$  for all  $i \neq j$
2. Using the pairwise distributions from step 1, compute the mutual information  $I(X_i; X_j)$  for each pair of nodes and assign it as the weight to the corresponding edge  $(X_i, X_j)$ .
3. Compute the maximum-weight spanning tree (MSWT):
  - a. Start from the empty tree over  $n$  variables
  - b. Insert the two largest-weight edges
  - c. Find the next largest-weight edge and add it to the tree if no cycle is formed; otherwise, discard the edge and repeat this step.
  - d. Repeat step (c) until  $n-1$  edges have been selected (a tree is constructed).
4. Select an arbitrary root node, and direct the edges outwards from the root.
5. Tree approximation  $P'(x)$  can be computed as a projection of  $P(x)$  on the resulting directed tree (using the product-form of  $P'(x)$ ).

# Summary:

## Learning and inference in BNs

- **Bayesian Networks** – graphical probabilistic models
- Efficient **representation** and **inference**
- Expert **knowledge** + learning from **data**
- Learning:
  - **parameters** (parameter estimation, EM)
  - **structure** (optimization w/ **score** functions – e.g., **MDL**)
  - **Complexity trade-off:**
    - NB, BNs and trees
- There is much more: causality, modeling time (DBNs, HMMs), approximate inference, on-line learning, active learning, etc.





# Online/print resources on BNs

## Conferences & Journals

UAI, ICML, AAAI, AISTAT, KDD

MLJ, DM&KD, JAIR, IEEE KDD, IJAR, IEEE PAMI

## Books and Papers

Bayesian Networks without Tears by Eugene Charniak. AI Magazine: Winter 1991.

Probabilistic Reasoning in Intelligent Systems by Judea Pearl. Morgan Kaufmann: 1998.

Probabilistic Reasoning in Expert Systems by Richard Neapolitan. Wiley: 1990.

CACM special issue on Real-world applications of BNs, March 1995

# Online/Print Resources on BNs

AUAI online: [www.auai.org](http://www.auai.org). Links to:

- Electronic proceedings for UAI conferences

- Other sites with information on BNs and reasoning under uncertainty

- Several tutorials and important articles

- Research groups & companies working in this area

- Other societies, mailing lists and conferences

# Publicly available s/w for BNs

List of BN software maintained by Russell Almond at  
[bayes.stat.washington.edu/almond/belief.html](http://bayes.stat.washington.edu/almond/belief.html)

several free packages: generally research only

commercial packages: most powerful (& expensive) is  
HUGIN; others include Netica and Dxxpress

we are working on developing a Java based BN toolkit here at  
Watson